

**MASARYK
UNIVERSITY**

FACULTY OF INFORMATICS

Source Code Quality impact on Pull Requests acceptance

Master's Thesis

ONDŘEJ KUHEJDA

Advisor: Assistant professor Bruno Rossi

Department of Computer Systems and Communications

Brno, Spring 2022

MUNI
FI

Declaration

Hereby I declare that this paper is my original authorial work, which I have worked out on my own. All sources, references, and literature used or excerpted during elaboration of this work are properly cited and listed in complete reference to the due source.

Ondřej Kuhejda

Advisor: Assistant professor Bruno Rossi

Acknowledgements

I would like to thank my supervisor, Bruno Rossi, for his guidance throughout the whole process.

Computational resources were supplied by the project “e-Infrastruktura CZ” supported by the Ministry of Education, Youth and Sports of the Czech Republic.

Abstract

TODO

Keywords

code quality, pull request, static code analysis

Contents

1	Introduction	1
1.1	Problem statement	1
2	Code quality in pull-based development	2
2.1	Code quality	3
2.2	TODO GitHub	4
3	Pull request acceptance	6
3.1	Repository level	6
3.2	Submitter level	7
3.3	Pull request level	7
3.4	TODO Code quality	8
3.5	TODO Unsorted	9
3.6	TODO Create table that compares already performed studies with my thesis	9
4	Data mining	10
4.1	GHTorrent database	11
4.2	git-contrast	12
4.3	TODO Projects selection	13
4.4	TODO Computational resources	14
5	Data analysis	15
5.1	TODO Scripts	15
5.2	Statistical methods	16
6	Evaluation	20
6.1	Python	20
6.2	Java	26
6.3	Kotlin	33
6.4	Haskell	39
6.5	C/C++	44
6.6	Programming languages and code quality impact . . .	50
6.7	Threats to validity	54

7	TODO Conclusion	56
7.1	TODO Comparison with the previous results	56
7.2	TODO Future work	56
	Appendix	57
	Bibliography	63

1 Introduction

1.1 Problem statement

- RQ₁** Which code issues are typically introduced by the pull requests?
- RQ₂** Are there some particular issues/code smells that affect the pull request acceptance?
- RQ₃** Is there a relationship between the source code quality and the pull request acceptance?
- RQ₄** Does code quality influence the time it takes to close a pull request?
- RQ₅** Is code quality impact higher in projects that are using some particular programming language?

2 Code quality in pull-based development

The pull-based development model created novel ways how can developers interact between each other. Instead of pushing code changes (patches) into one central repository, developers can work in more decentralized and distributed way. This is mainly done by using distributed version control systems such as Git. Git enables developers to clone repositories and thus to work independently on projects. Furthermore, the Git's branching model helps developers to keep track of repository changes and helps to handle the conflicts between the different changes of the same code base.

To furthermore ease the complicated process of resolving conflicts between different changes (of the same code base) and to provide a more user-friendly environment for developers, platforms such as GitHub was created. These platforms adds new ways how the developers can interact beyond the basic functionality of Git:

- The forks enables to create the server-side copy of the repository.
- Pull requests (on some platforms called merge requests) enables to merge code directly on the platform.
- Users can report issues found in the projects; therefore, platform can also serve as a bug-tracking system.
- The comments can be added to the pull requests and issues in order to build up social interaction between developers.
- Users can star projects and follow other users, projects, pull requests or issues.

In this study, I choose to use GitHub as the main source for data mining. GitHub is one of the leading platforms that enables pull-based collaboration between developers. GitHub hosts huge amount of publicly available repositories and GitHub also provides public REST API that can be easily leveraged for data mining.

The aim of this thesis is to obtain large amount of data about GitHub projects and analyze the pull request in regard of their code quality. How the code quality can be analyzed and how the GitHub

platforms contributes to quality of the code itself is discussed in the following chapters.

- **TODO** cite: An Exploratory Study of the Pull-based Software Development Model

2.1 Code quality

Code quality is very important aspect of every program — software with high code quality has competitive advantage, is more stable and is also more maintainable then software which is poorly written.

To be able to evaluate the software in regard of its quality, there needs to be some way how can be code quality measured. The testing can be used exactly for this purpose — as a tool for measuring the quality of the source code. There are multiple ways how can be testing performed. Testing techniques can be divided into two categories: static and dynamic testing techniques.

In order to use dynamic testing techniques on large number of programs, there are two large obstacles — the program needs to be executed and there needs to be some inputs (with expected outputs) that can be then used for testing. Program execution can be problematic. Some programs needs to be compiled before they can be executed; others requires special environment for its execution (specific hardware, operating system or shared libraries required by the program). Moreover, the most of the programs does not have sets of input that can be used for testing. There exists some techniques that can be used also without the predefined inputs such as fuzzing, but these techniques are usually time-consuming. Because of that, dynamic testing techniques are not viable option when dealing with the large number of programs.

On the other hand, static testing methods suits the analysis of the large number of programs better. Static techniques include usage formal and informal reviews, walkthroughs and inspections; however, these techniques are performed by humans and therefore are not usable for large datasets. Because of that, in this thesis, the quality of the given source code is evaluated using the tools for automatic static analysis (called linters). Linters are used to find defects and code smells in the source code without the need of source code's execution.

There are several categories of issues which can be detected using linters. Source code can be checked if it follows a conventions of the given programming language. For instance, Python has an official style guide for Python code — PEP 8¹. This guide defines the conventions that should be followed such as proper indentation of the code blocks, maximum line length or naming conventions.

Furthermore, code can be analyzed against refactoring related checks; for instance linter can detect if some part of the code is redundant and therefore could be omitted. Linters can also detect actual errors such as type mismatches or syntax errors.

However, it is important to note that not all linters have the same capabilities. Number of issues which can be detected by the given linter also heavily depends on the programming language of the studied source code. Which linters were used for the purposes of this thesis is discussed later in the text.

- **TODO** cite: https://www.utcluj.ro/media/page_document/78/Foundations%20of%20software%20testing%20-%20ISTQB%20Certification.pdf

2.2 TODO GitHub

- GitHub issues and code quality
- Ways to merge code
 - An Exploratory Study of the Pull-based Software Development Model
- PRs and code review
- PRs CI/CD and code quality
 - Wait for It: Determinants of Pull Request Evaluation Latency on GitHub [1]
 - * CI and latency

1. <https://www.python.org/dev/peps/pep-0008/>

- Trautsch et al. [2] analyzed several open-source projects in regards to usage of static analysis tools. They found out that incorporating a static analysis tool in a build process reduces the defect density.

3 Pull request acceptance

Pull request acceptance is a problem that has been studied multiple times. Several surveys were performed in order to understand why pull requests are being rejected.

Gousios et al. [3] surveyed hundreds of integrators to find out their reasons behind the PR rejection. Code quality was stated as the main reason by most of the integrators; code style was in the second place. Factors that integrators examine the most when evaluating the code quality are style conformance and test coverage.

Kononenko et al. [4] performed a study of an open-source project called *Shopify*; they manually analyzed PR's and also surveyed *Shopify* developers. They found out that developers associate the quality of PR with the quality of its description and with the revertability and complexity of the PR.

The reasons why contributors abandon their PRs were also studied [5]. The reason number one was the "Lack of answers from integrators."; moreover, the "Lack of time" and the "Pull request is obsolete" was also often stated as the main reason.

Even though the different open-source communities solve the problem of pull request acceptance in a different manner, three main governance styles can be identified — protective, equitable, lenient. Protective governance style values trust in the contributor-maintainer relationship. The equitable governance style tries to be unbiased towards the contributors, and the lenient style prioritizes the growth and openness of the community [6]. Each style focuses on different aspects of PR. Tsay et al. [7] identified the following levels of social and technical factors that influence the acceptance of the PR — *repository level*, *submitter level*, and the *pull request level*.

3.1 Repository level

The *repository level* is interested in the aspects of the repository itself, such as the repository age, number of collaborators, or number of stars on the GitHub.

For instance, the programming language used in the project also influences the acceptance of the PRs. Pull requests containing Java,

JavaScript, or C++ code have a smaller chance to be accepted than PRs containing the code written in Go or Scala [8].

Furthermore, older projects and projects with a large team have a significantly lower acceptance rate [7].

The popularity of the project also influences the acceptance rate — projects with more stars have more rejected PRs [7].

3.2 Submitter level

The *submitter level* is concerned about the submitter's status in the general community and his status in the project itself. There are several parameters that can be considered when evaluating the submitter's status.

PRs of submitters with higher social connection to the project have a higher probability of being accepted [7].

Submitter status in the general community plays an important role in PR acceptance. If the submitter is also a project collaborator, the likelihood that the PR will be accepted increases by 63.3% [7].

Moreover, users that contributed to a larger number of projects have a higher chance that their PR will be accepted [9].

The gender of the submitter is another factor that plays a role in PR acceptance. A study showed that woman's PR are accepted more often, but only when they are not identifiable as a woman [10].

Personality traits also influence PR acceptance. The *IBM Watson Personality Insights* were used to obtain the personality traits of the PR submitters by analyzing the user's comments. These traits were then used to study PR acceptance. It has been shown that conscientiousness, neuroticism, and extroversion are traits that have positive effects on PR acceptance. The chance that PR will be accepted is also higher when the submitter and closer have different personalities [11].

3.3 Pull request level

The *pull request level* is interested in the data that are connected to the PR itself. For instance, on the *PR level*, one can study if there is a correlation between PR acceptance and the number of GitHub

comments in the PR. Another parameter that can be used is “Number of Files Changed” or “Number of Commits”.

One of the factors that negatively influence the acceptance rate is the already mentioned number of commits in the pull request. The high number of commits decreases the probability of acceptance. On the other hand, PR’s with only one commit are exceptions — they have a smaller chance to be accepted than pull requests which contain two commits [9].

Another observation is that more discussed PR’s has a smaller chance to be accepted [7]. Another study did not find a large difference between accepted and rejected PR’s based on the number of comments but found that discussions in rejected PR’s have a longer duration [12].

Proper testing is the crucial part of every project, and therefore it also influences the pull request acceptance. PR’s including more tests have a higher chance to be accepted, and an increasing number of changed lines decreases the likelihood of PR acceptance [7].

Testing plays a significant role in discovering bugs and therefore leads to higher code quality. On the other hand, many test cases do not have to mean that code has a high quality. The code quality is an essential factor on the *pull request level*, therefore, is this study’s main interest. Works that are also interested in the code quality and the pull request acceptance are examined in the following chapter.

Another factor that is closely tied to code quality is the code style. This factor has a small (but not negligible) negative effect on acceptance. This means that PRs with larger code style inconsistency (with the codebase) have a smaller chance of being accepted [13].

3.4 TODO Code quality

Although most integrators view code quality as the most important factor regarding PR acceptance, to the best of my knowledge, only one study was performed to discover whether there is a connection between the PR’s acceptance and its quality.

- Does code quality affect pull request acceptance? [14]

3.5 TODO Unsorted

- study “Influence of Social and Technical Factors” [7] was replicated [11]
- Replication Can Improve Prior Results: A GitHub Study of Pull Request Acceptance [15]
 - contains interesting table with factors that influences acceptance
- Pull Request Decision Explained: An Empirical Overview [16]
 - also contains interesting table with factors that influences acceptance
- An Exploratory Study of the Pull-Based Software Development Model [17]
- Which Pull Requests Get Accepted and Why? A study of popular NPM Packages [18]
- Rejection Factors of Pull Requests Filed by Core Team Developers in Software Projects with High Acceptance Rates [19]
- Pull Request Prioritization Algorithm based on Acceptance and Response Probability [20]

3.6 TODO Create table that compares already performed studies with my thesis

4 Data mining

TODO: update graph Information about the pull requests are retrieved




Figure 4.1: The `pr_quality.py` workflow

using the `pr_quality.py` script. This script takes names of the projects that will be analyzed as the input and it outputs the JSON files containing the requested data. This script uses internally two other scripts — `gh_info.py` and `git-contrast`.

`gh_info.py` is responsible for querying the GHTorrent database in order to obtain data about the projects. The GHTorrent database is an offline mirror of data offered through the Github REST API. `gh_info.py` returns a JSON file with the information about the project such as number of stars, number of contributors or information about pull requests and their commits.

However, the Github REST API lacks the information about the code quality of the pull requests. This is where the `git-contrast` comes into the play. `git-contrast` is the command-line application which analyzes the code quality of the given pull request using the

external linters. This application is further discussed in the following sections.

TODO: mention that data from REST API are not complete (GH API limit)

4.1 GHTorrent database

As stated before, the script called `gh_info.py` uses the GHTorrent database in order to retrieve GitHub data. GitHub REST API can be leveraged to obtain many interesting factors which can possibly influence the acceptance of pull requests. All the data that are obtained using the `gh_info.py` are listed in the following table:

Table 4.1: Data retrieved from the GHTorrent

Level	Variable	Factor
Repository level	Project name	✗
	Programming language	✓
	Time of creation	✓
	Number of forks	✓
	Number of commits	✓
	Number of project members	✓
	Number of watchers	✓
Submitter level	Username	✗
	Number of followers	✓
	Status in the project	✓
Pull request level	Pull request ID	✗
	Is PR accepted?	✓
	Time opened	✓
	Head repository	✗
	Head commit	✗
	Base commit	✗
	Number of commits	✓
	Number of comments	✓

Variables marked with ticks (✓) are factors that can possibly influence code quality and they can be used for pull request acceptance analysis. Other variables (✗) are not meant to be used as an part of an

data analysis itself, but are kept here for better orientation; and some of them are later used by the `git-contrast` tool (in order to pull the commits which will be subsequently analyzed by linters).

4.2 `git-contrast`

`git-contrast` is the command line application that I implemented in order to be able to analyze the code quality of the given pull request. `git-contrast` expects two commit hashes on the input and returns the information about the change in code quality between these commits on the output. This is done by running the linter on the files in the state of the first commit and then in the state of the second commit. The number of found code quality issues is then written to the standard output.

To measure the change of the quality in the pull request, we simple run the `git-contrast` on the “head commit” and the “base commit” of the given pull request. `git-contrast` supports several linters; which linter will be used is determined by the file extension of the tested file. Linters that are supported by `git-contrast` are listed in the following table:

Table 4.2: Linters supported by the `git-contrast`

Linters	Programming languages	File extensions
OCLint	C/C++	.c, .cpp and .h
HLint	Haskell	.hs
ktlint	Kotlin	.kt and .kts
PMD	Java	.java
Pylint	Python	.py

The most problematic was to statically analyze the C/C++ source files because some linters also need the information how the source code should be compiled. Luckily, this information can be usually automatically obtained from the makefiles. Another problem is the speed. At first, I was using the Cppcheck linter for the static analysis of C/C++ but I was forced to switch to the OCLint in order to shrink the total execution time of the static analysis.

TODO: add information about versions of linters, issue categories etc.

4.3 TODO Projects selection

Criteria (data from 2019-06-01):

- is in the top 150 most favorite projects written in the given language
- 200+ pull requests and less than 5000
- PRs merged using GitHub
- <https://github.com/EvanLi/Github-Ranking>
- at least 85 % of files are source files written in the given language
- project is a program or program collection (not a book with the script etc.)
- The popularity of the project influences the acceptance [7].
- <https://dl.acm.org/doi/abs/10.1145/2597073.2597122>
- <https://dl.acm.org/doi/abs/10.1145/3379597.3387489>
- <https://zenodo.org/record/3858046>
- <https://github.com/XLipcak/rev-rec>
- <https://ghtorrent.org/>
 - <https://github.com/gousiosg/pullreqs>
 - How can I cite this work? (on the web)
- Kalliamvakou et al. noted that data about PR's mined from GitHub are not always reliable, because PR can be also merged using several different approaches.
 - <https://dl.acm.org/doi/10.1145/2597073.2597074>
 - [17]

4.4 TODO Computational resources

5 Data analysis

5.1 TODO Scripts

- Put this into appendix?

Data preprocessing

In order to simplify analysis of retrieved data, I created the script that takes multiple JSON files with the data about each individual project and converts them into the CSV files. Each row in the CSV file represents some pull request. This script also filters the pull requests which are not suitable for the analysis — PRs that do not contain any source code written in the primary language or PRs that contained corrupted files (the linter was unable to analyze those files).

Pull request classification

The retrieved data about the pull request were subsequently analyzed in order to answer my research questions. For this analysis was used the Python script¹ provided by Lenurdazzi et al. (this script was used for machine learning classification methods)

Pull request regression

Main script

- rest of the analysis
- divided by research questions
- aggregation of results from the classification and regression
- export to \LaTeX and PDF

1. <https://figshare.com/s/d47b6f238b5c92430dd7>

5.2 Statistical methods

Research questions 1 to 4 were analyzed separately for each programming language.

TODO: For each test, state the rationale behind the test and assumptions about data.

- <https://www.scribbr.com/statistics/statistical-tests/>

Which code issues are typically introduced by the pull requests?

At first, in order to answer the **RQ₁**, I summarized the retrieved data for each project — I counted how many suitable pull requests were analyzed and how many of them are accepted/rejected. Then I created a scatter plot between number of stars and percentage of accepted PRs.

I also summarized all pull requests regardless of their project. I computed average number of introduced issues, fixed issues etc. Then I created heat map that shows how many PRs introduced/fixed some specific number of issues.

Then for each issue individually I computed how many accepted/rejected pull request introduced/fixed this issue, how many times this issue occurred in some pull request etc. I created multiple lists of issues sorted by some parameter. I sorted issues by the number of rejected/accepted PRs that fixed/introduced them. I also listed issues and percentage of PRs that changed their quality. I examined the issues that were fixed in larger number of PRs then introduced. Then I created a scatter plot that shows which issue category contains the most common issues.

These steps were applied individually for each programming language to determine how does the average PR look line in terms of code quality.

Are there some particular issues/code smells that affect the pull request acceptance?

In order to discover issues that affects the acceptance of pull requests most, the classification models were created. The aim of these models is to classify pull request into two groups (accepted PRs and rejected

PRs) by using the information about quality change in the given pull request. Multiple classification algorithms were used:

LogisticRegression TODO

RandomForest TODO

GradientBoost TODO

ExtraTrees TODO

DecisionTrees TODO

Bagging TODO

AdaBoost TODO

XGBoost TODO

Each of those algorithms was run on three different datasets — dataset with quality change, dataset containing only introduced issues and dataset with only fixed issues. In the first dataset, the quality change for some issue was represented by the integer and this integer was negative if the issue was fixed in the PR and positive if the issues was introduced. The other datasets were created by filtering positive/negative values from the first dataset.

In order to recognize issues that have some effect on the PR acceptance, the *drop-column importance* mechanism² was used.

- **TODO:** reference the paper
- **TODO:** correlation matrix
- **TODO:** reference/explain the model reliability (from the appendix)

2. <https://explained.ai/rf-importance/>

Is there a relationship between the source code quality and the pull request acceptance?

RQ₃:

- classification
- PCA scatterplot
- contingency matrices
- contingency matrices for PR's that contain only modified source code files vs for all of them
- contingency matrices separately for each project
 - According to Cochran (1952, 1954), all expected counts should be 10 or greater. If < 10 , but ≥ 5 , Yates' Correction for continuity should be applied.
- ROC curves and AUCs
- mention that p-values were not adjusted

Pull request that are adding or removing some files has large influence on code quality. If the number of removed/added files has large impact on PR acceptance, then it can be a large threat to validity of the independence test. The pull request acceptance can be also influenced by quality of files which were not linted (was written in non-primary language). To eliminate the risk that the test was influenced, same test was performed on pull requests that only modified some source files and these files were written in the primary language.

Does code quality influence the time it takes to close a pull request?

In order to find the possible link between the code quality and time it takes to close a PR, the regression algorithms were used. At first, the dataset was split into two parts — training and test set. After that, the regression model was trained on the training set. Then, the importance of individual quality issues was determined using the *permutation importance* mechanism. To evaluate the model itself, the

model was used to predict the time based on the data from test set. Metrics such as *mean absolute error* (MAE), *mean squared error* (MSE) and *coefficient of determination* (R^2) were computed using the predicted and expected values.

- **TODO:** correlation matrix

Is code quality impact higher in projects that are using some particular programming language?

RQ₅:

- Compare results from previous steps.
 - Statistically compare the parameters obtained for each programming language. (check if the pull requests from different languages and retrieved parameters follow the same distributions)
- What is the effect of the programming language on the acceptance and time it takes to close a pull request?
 - ANOVA
 - classification (machine learning)

6 Evaluation

The following chapter is dedicated to the findings from my research. The first five subchapters focus on individual programming languages — here I am giving the answers to the first four research questions. The last research question (**RQ5**) is answered afterward. At the end of this chapter, I am discussing possible threats to validity that could eventually influence the outcomes of my study.

6.1 Python

In order to analyze the influence of code quality on the pull request acceptance, 20 projects from the Python ecosystem were selected. In total, 9452 pull requests were analyzed, and 73 % of these PRs were accepted. Pull requests were more accepted in less popular projects, as can be seen in the following scatterplot:




Figure 6.1: Stars and pull request acceptance

On average, one pull request introduced 5.36 issues and fixed 2.44 issues; an accepted pull request introduced 4.62 and fixed 1.99 issues, and rejected pull request introduced 7.86 issues and fixed 4.43 on average. 5 % trimmed mean was used to compute these values.




Figure 6.2: Pull requests and quality

In the analyzed pull requests, Pylint detected 222 different issues. The conventions dominated the list of issues that were fixed/introduced in the largest number of pull requests. The convention that was fixed/introduced in the largest number of pull requests is missing-

function-docstring (in 37 % of PRs); conventions `invalid-name`, `line-too-long` and `consider-using-f-string` were fixed/introduced in over 20 % of pull requests. There were 15 issues that were fixed/introduced in more than 10 % of PRs, and 72 issues were in over 1 % of PRs (out of the 222 issues which were found in the pull requests). There were nine issues that were present in the analyzed pull requests but did not influence their quality (the number of these issues was not changed by any pull request). 13 issues were introduced/fixed in only one pull request, and 10 of them are issues classified as errors. The most common error is `import-error` (24 % of PRs); however, I suspect that there will be many false positives that arise due to linting in the isolated environment. Sixty issues were fixed in more PRs than they were introduced. They are 24 more PRs that fixed the warning `super-init-not-called` than the PRs that introduced it.




Figure 6.3: Pylint issues and % of PRs which fixed/introduced them

The most important Pylint issue in regards to the PR acceptance is the `syntax-error`. XGBoost classifier gives this error the 1.2 % importance. However, other classifiers consider this error less important. On average importance of the `syntax-error` is only 0.3 %. The `syntax-error` was introduced in 17 projects. On average, rejected pull request introduced 0.027 `syntax-error`s, and the average accepted pull request even fixed 0.001 `syntax-error`s.

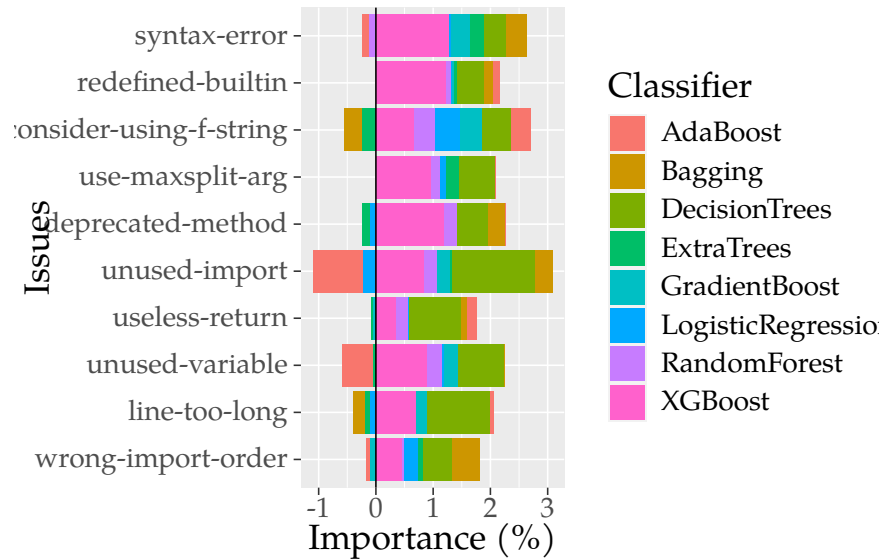


Figure 6.4: Ten most important Pylint issues

When only introduced issues were considered, the list of the most important issues looked differently. On the other hand, there are some issues that appeared in the top 10 in both lists: `syntax-error`, `unused-variable` and `unused-import`. The `syntax-error` is considered the most important issue by both methods.

When only the information about fixed issues is used, the most important issue is `f-string-without-interpolation` (in terms of acceptance). However, no classifier gives this issue importance over one percent.

In order to visualize the difference in quality between accepted and rejected PRs, I created PCA scatter plot:




Figure 6.5: PCA scatter plot

In the PCA scatter plot, there is no visible difference between rejected and accepted pull requests.

To understand if the presence of some issue in the PR influences its acceptance, I created contingency matrices and performed a chi-square test of independence:




Figure 6.6: Relation between presence of issue and PR acceptance

As can be seen in Figure 6.6, the observed number of rejected pull requests which contained some defected is higher than expected. For chi-square test, $p < 2.2 \times 10^{-16}$ and therefore, the hypothesis that presence of some issue and PR acceptance are independent is rejected on significance level $\alpha = 0.05$. However, the Cramer's $\phi_c \approx 0.092$; therefore, the association between issue presence and acceptance is weak. This conclusion also supports the fact that AUROC for trained classification models is only slightly over 0.5. The average AUC for all models is 0.534.

When considering only PRs that solely modified some source files, $p < 5.548 \times 10^{-10}$ and therefore also here the presence of some code quality issue in the PR influences the PR acceptance. Similar to the previous test, the $\phi_c \approx 0.087$; therefore, the association between the presence of the same issue and PR acceptance is weak.

Almost identical results were obtained when the chi-square test was performed separately for each issue category.

When the projects were considered individually, only for nine of them the $p < \alpha$. In these projects, the poor code quality had a negative impact on PR acceptance. In the rest of the projects, the presence of

some code quality issue does not seem to have an effect on the PR acceptance.

The quality of the code does not seem to have an effect on the time it takes to close a pull request. All of the trained regression models have a negative R^2 score (when evaluated on the test set). This means that trained models are worse at predicting the time than a constant (mean value). Similar results were obtained when only introduced issues were considered and also when only fixed issues were considered.

6.2 Java

The next programming language that was analyzed is Java. In total, the 8887 pull requests were linted, and 73 % of these pull requests were accepted. On average, the one pull request introduced 20 new PMD issues but, at the same time, also fixed 18 other issues.

Like in the Python projects, the pull request from the less popular project were more likely to be accepted than pull requests from more popular projects.




Figure 6.7: Stars and pull request acceptance

Only 1366 pull requests (from the total of 8887 pull requests) did not change the quality of the source code (did not fix nor introduce some PMD issues). The PMD linter was able to detect 253 different issues in the given pull requests. Most of the introduced issues were issues related to the code style. In total, all of the pull requests introduced over a million code-style issues.




Figure 6.8: Pull requests and quality

The issue that was introduced in the largest number of pull requests is `CommentRequired` (documentation issue). Another frequent issues are `LocalVariableCouldBeFinal`, `MethodArgumentCouldBeFinal` (code style issues) and `LawOfDemeter` (issue in code design). These issues are the only issues that were introduced in more than 3000 pull requests. Similarly, the list of issues that were fixed in the

largest number of the pull request is dominated by the very same issues.




Figure 6.9: PMD issues and % of PRs which fixed/introduced them

As can be seen in Figure 6.9, the documentation issues tend to appear in a large number of pull requests (24 % on average). Moreover, the typical code style issue appeared in 11 % of pull requests. On the other end of the spectrum, an average issue indicating an error-prone construct is present in only two percent of pull requests.

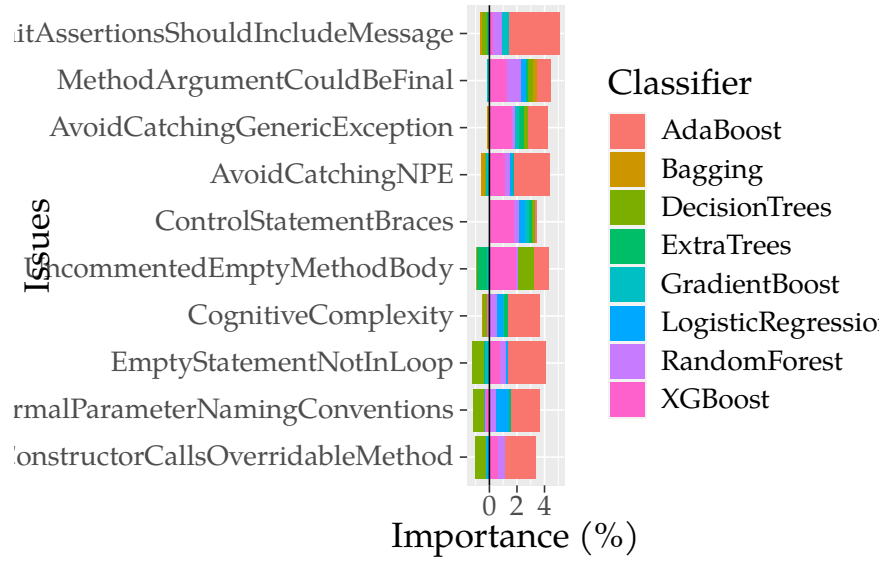


Figure 6.10: Ten most important PMD issues

The most important PMD issue is `JUnitAssertionsShouldIncludeMessage`. The average importance of this issue is only 0.6 %. However, the AdaBoost classifier gives this issue 3.7 % importance. The 0.89 issues of this type are introduced in an average accepted pull request. I suspect that the pull requests that are adding a larger number of tests to the codebase have a higher probability of being accepted. At the same time, these pull requests also have a higher probability of introducing the `JUnitAssertionsShouldIncludeMessage`. This can be the reason why this issue has the largest importance. This also supports the study that shows that the acceptance likelihood is increased by 17.1 % when tests are included [7]. However, another performed study indicates that the presence of test code does not influence PR acceptance [17].

The PCA scatter plot for Java pull requests looks as follows:




Figure 6.11: PCA scatter plot

In the PCA scatter plot, there is no visible difference between rejected/accepted pull requests.

To understand the relationship between acceptance and the introduction of a quality issue, the chi-square test was performed.




Figure 6.12: Relation between presence of issue and PR acceptance

The $p = 9.132 \times 10^{-14} < \alpha$ and $\phi_c = 0.079$; therefore, there is a weak relation between acceptance and issue presence. Similar results were obtained when only PRs that solely modified the source code of the main language were considered and also when the test was performed individually for each issue category.

17 out of the 20 Java projects contained a sufficient number of pull requests to perform the chi-square tests. In nine of them, the code quality and acceptance are not independent. Unexpectedly, in one of the projects (alibaba/fastjson) the presence of an issue has a small positive effect on the acceptance.

The PMD issues seem to have some effect on the time it takes to close a pull request when considering only R^2 computed for each model. However, the R^2 value is usually not a good metric for evaluate non-linear models; it can reveal some information about the model, but it does not give us information on how accurate the model is. There are three models that have $R^2 > 0.4$: Bagging, GradientBoost and RandomForest. The linear regression has $R^2 = 0.1257$; therefore for this model, 13 % of the variance in time to close a PR can be explained by quality issues. However, all of the models have high mean absolute error (MAE). The average MAE value for all of the models is

3934338 \approx 46 days and 87 % of all analyzed Java pull requests were closed within one month. Therefore these models are basically useless in practice. The other models (when considering only rejected/fixed issues) yielded similar results. To conclude, the found quality issues do not seem to have an effect on the time to close a pull request.

6.3 Kotlin

The 20 projects were also selected from the Kotlin ecosystem. The average analyzed pull request was from a project that has ten thousand stars and introduced nine issues and fixed only four. The 7514 pull requests were analyzed (using the *ktlint* linter), and 80 % of them were accepted. The trend that maintainers of popular projects reject more pull requests can also be observed in the Kotlin community.




Figure 6.13: Stars and pull request acceptance

Only 20 different issues were detected by the *ktlint* in the analyzed projects; however, this is expected since the *ktlint* is focused only on a small set of quality issues.

The `indent` is the issue that was introduced in the largest number of pull requests (2598). It is the only issue that was introduced in more than a thousand pull requests. It is also the issue that was fixed in the largest number of pull requests. The official Kotlin convention is to use the four spaces for indentation¹, and the `indent` issue signifies that this convention was violated. This issue influenced the code quality of more than half of the pull requests. However, this can be caused by projects whose standards do not follow the official recommendations.

Other often violated *ktlint* rules are `no-wildcard-imports`, `final-newline`, and `import-ordering`. On the other end of the spectrum, the rule `no-line-break-after-else` was violated only once.

1. <https://kotlinlang.org/docs/coding-conventions.html>




Figure 6.14: Pull requests and quality

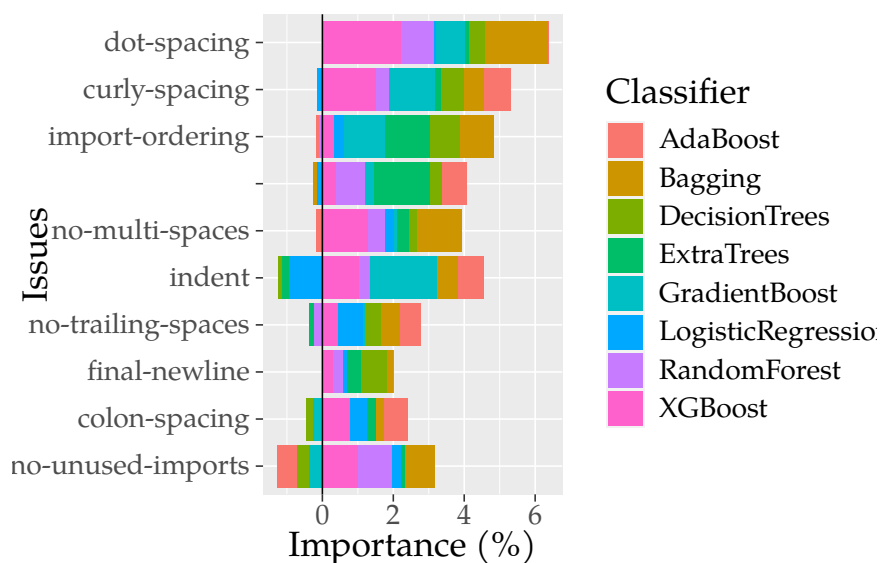


Figure 6.15: Ten most important ktlint issues

The issue with the highest importance average is dot-spacing. The Bagging classifier gives 1.7 % importance to this issue. The importance obtained from other classifiers is smaller — the average importance is 0.8 %. However, this issue was introduced only in 18 PRs (13 times in the rejected pull request). Furthermore, seven accepted and seven rejected pull requests fixed this issue. Therefore the impact of this issue is disputable.

It is worth mentioning that fourth most important issue does not have a name (given by *ktlint*). This issue usually indicates an invalid Kotlin file. This issue has high importance (relative to the other issues) also when the only fixed and also when only introduced issues were taken into account during the classification. This issue was introduced by 90 rejected PRs and by 51 accepted PRs.

When using only introduced issues, the most important issue is indent. This issue is also most important when only the fixed issues are considered. As being said before, in projects that are using non-standard indentation, this issue is a false positive.

The PCA scatter plot was also created for the Kotlin programming language. The first principal component explains almost all variance

in the code quality of pull requests. However, the difference between rejected/accepted pull requests is not apparent from the PCA plot:




Figure 6.16: PCA scatter plot

To understand the link between acceptance and the introduction of some quality issue, I performed the chi-square test on Kotlin dataset. The $p < 2.2 \times 10^{-16}$ and $\phi_c \approx 0.095$; therefore, the presence of some issue has a small negative effect on acceptance (similarly to the Java and Python). Furthermore, three classifiers (*Bagging*, *GradientBoost*, and *RandomForest*) have AUC for ROC curve above 60, and the average AUC is 57.58.




Figure 6.17: Relation between presence of issue and PR acceptance

However, taking into account solely the PRs that only modified some source code, the $p = 0.627$, thus the acceptance and issue presence are independent (in this context).

Only 12 of the projects have a sufficient number of pull requests to evaluate the chi-square test. There are four projects where the presence of some issue has a small impact on the PR acceptance (the average Cramer's V is $\phi_c = 0.18$).

To analyze the relation between the code quality and time that is required to close a PR, I applied several regression techniques also to the Kotlin dataset. For linear regression, $R^2 = 0.164$, therefore the trained model is able to explain 16 % of the variance in the time to close a PR. The $MAE = 2375121 \approx 27$ days; therefore, the model does not perform so well on the dataset, taking into consideration that 89 % of pull requests were closed within one month. The mean absolute error for other models was similar to the MAE obtained for linear regression.

6.4 Haskell

Haskell is the only purely functional programming language that was analyzed. The 18 out of 20 selected Haskell projects have under the 5000 stars. There are only two exceptions: PureScript with 7632 stars and Pandoc, which has over 15000 stars. The Pandoc has the also smallest percentage of accepted pull requests. However, excluding the Pandoc, there is no visible connection between the number of stars and acceptance in the selected projects. When the outliers are filtered, the trend tends to be the opposite of previous languages: more accepted are pull requests of projects with more stars. However, only 20 projects are not sufficient to make such conclusions about the whole population of Haskell projects.




Figure 6.18: Stars and pull request acceptance

The 6949 pull requests were analyzed. Interestingly, in over 60 % of pull requests, no change in the code quality was detected. Moreover, the *HLint* is able to recognize a large number of different issues (321 issue types were detected in selected pull requests). On the other

hand, some issues were counted twice because they appeared as a suggestion but also as a warning (in the different contexts). These facts can indicate that a large number of submitted pull requests follow high-quality standards.




Figure 6.19: Pull requests and quality

Seventy-eight percent of pull requests were accepted, and the average pull request introduced only 0.6 issues and fixed 0.3 issues. The

most common types of issues were suggestions and warnings. The error that was introduced in the largest number of pull requests is `Use-newTVarIO`, and this error was introduced only in 8 pull requests. The most common suggestions were `Redundant-bracket` (introduced in 499 PRs) and `Redundant-$` (444 PRs). The warning `Unused-LANGUAGE-pragma` was introduced in 323 pull requests and `Eta-reduce` warning in 214 of them. There were only ten issues that were introduced in 100 and more pull requests; and another 105 issue types were detected in the analyzed code, but no PR introduced any of those issues.




Figure 6.20: HLint issues and % of PRs which fixed/introduced them

The most important Haskell issue is the suggestion `Use-if`. However, no classifier gives this issue importance over one percent. Therefore the actual impact of this issue is disputable. This issue was introduced in 18 rejected PRs and fixed in 11. There are 19 accepted PRs that introduced `Use-if` and 27 accepted PRs that fixed it.




Figure 6.21: Ten most important HLint issues

When only introduced issues were taken into account, the most important issue is `Move-brackets-to-avoid-$` (suggestion). The AdaBoost classifier gives this issue 1 % importance, although the average importance is only 0.4 %.

In the context of fixed issues, the most important is warning `Use-fewer-imports` with average importance again only about 0.4 %.

The PCA scatter plot was also generated for the Haskell language. Similar to the results in already analyzed languages, there is no apparent difference between accepted and rejected pull requests.




Figure 6.22: PCA scatter plot

For the chi-square test, the $p = 0.001438 < \alpha = 0.05$ and Cramer's V is only $\phi_c = 0.038$; therefore, the presence of an issue in the PRs has only a small negative impact on the acceptance of the pull request. Similar results were obtained when only the pull requests that contain exclusively some modified code were considered. Furthermore, tests for the individual issue types also yielded similar results. Unfortunately, there is only a small number of pull requests that introduced some errors; therefore the chi-square test cannot be performed on this issue category. The average AUC computed for ROC curves is around 50 — the classification algorithms were unable to distinguish between the accepted and rejected PRs using the code quality.




Figure 6.23: Relation between presence of issue and PR acceptance

The 13 projects contain a sufficient number of pull requests; the acceptance and the issue presence are not independent only in four of them (there, the issue presence have a small negative impact on the acceptance). For the `haskell/aeson` project, the Cramer’s V is 0.282 — the association is “medium”.

The issues detected by *HLint* do not seem to have an impact on the time it takes to close a pull request. All trained models have negative R^2 . When only fixed issues were used for regression, there were three models with positive R^2 : Bagging (0.0315), ElasticNet (0.0085), and RandomForest (0.0229). However, all of them have high mean absolute error: Bagging ($2193658 \approx 25$ days), ElasticNet (2255678), and RandomForest (2201347).

6.5 C/C++

The C and C++ programming languages are analyzed together because they share a lot of similarities. This usually enables use of the same linter for both languages. Moreover, it is not uncommon that projects that are written in C++ also contain some C code and vice

versa. The nine selected projects have more code written in C, while the rest of the 11 projects is more C++-oriented.

In analyzed projects, there is no visible connection between the acceptance and the number of stars.




Figure 6.24: Stars and pull request acceptance

I analyzed 8774 C/C++ pull requests. Seventy-seven percent of them have been accepted. The typical pull request introduces 0.25 issues and fixes 0.12 issues; the typical rejected PR introduces 0.79 issues, and the typical accepted PR only 0.15 issues. The 79 % of pull requests did not change the quality of the source code (in terms of the *flawfinder* quality rules).




Figure 6.25: Pull requests and quality

The most common type of issue is the note. The least common are errors. The *flawfinder* was able to identify 137 different issues in the studied PRs. All of the top ten issues (in terms of number of PRs which introduced them) are notes. The most common note is buffer-char (“Statically-sized arrays can be improperly restricted leading to potential overflows or other issues...”). The most common

error is `buffer-strcat` (“Does not check for buffer overflows when concatenating to destination...”), and it is the 11 most introduced issue (introduced in 69 pull requests). There are 36 issues that were present in the analyzed code, but they were not introduced in any pull request; 21 of them are errors.




Figure 6.26: flawfinder issues and % of PRs which fixed/introduced them

Classification algorithms rank as the most important issue the `format-printf` (“If format strings can be influenced by an attacker, they can be exploited...”). However, this issue is only a *note*. Therefore it does not have to indicate a defect (there will probably be a large number of false positives). AdaBoost and XGBoost algorithms give this issue importance of 1 %. The average importance is 0.7%. This issue is also most important when only introduced issues are considered.

The second most important issue has average importance of only 0.26 %.

The most important error is `buffer-StrCpyNA` (“Does not check for buffer overflows when copying to destination...”) with average importance of only 0.9 %. This error is the sixth most important issue.




Figure 6.27: Ten most important flawfinder issues

When considering only fixed issues, the `buffer-read` is the most important issue (note); however, the average importance is only 0.28 %.

The PCA analysis does not reveal any significant difference between the accepted and rejected pull requests (in terms of code quality).




Figure 6.28: PCA scatter plot

Based on the chi-square test, the presence of an issue in the PR has a small negative impact on the PR acceptance ($\phi_c = 0.117$). However, When considering only pull requests that solely modified some source files, Cramer's V $\phi_c = 0.024$ and $p = 0.1 > \alpha$ — in this settings, the issue presence does not influence acceptance.




Figure 6.29: Relation between presence of issue and PR acceptance

Some small impact was discovered when the chi-square test was performed separately for each issue category (the $p < 2.2 \times 10^{-16}$ and $\phi_c \approx 0.1$ for each category). Furthermore, in 6 out of 11 projects which have enough data to perform and evaluate the chi-square test, the presence of some issue in the PR has a negative effect on the PR acceptance. In the `minetest/minetest` and `pybind/pybind11` projects, this effect is moderate; for other projects, the association is small.

In the case of C/C++, the time to close a pull request seems not to be related to found issues. All the models have negative R^2 , except the ElasticNet regressor. For the ElasticNet, $MAE = 4681624$ (the mean absolute error is 54 days) — therefore, this model also cannot be used to predict the time to close a PR. Models considering only rejected issues and also models considering only accepted issues have yielded similar results.

6.6 Programming languages and code quality impact

Comparing the code quality of projects written in different programming languages is a difficult task. Each language has different programming constructs, syntax, and type system. For instance, Python,

which is a dynamically-typed multi-paradigm programming language, has completely distinct characteristics from Haskell, which is a purely functional programming language with a strong, static type system.

Moreover, every linter is different and has a unique set of rules. The *ktlint* is focused on code clarity and community conventions, whereas *flawfinder* checks code for potentially dangerous functions. On the other hand, the *PMD* is a more general-oriented linter that contains a large set of rules for the Java programming language. Lastly, the *HLint* is oriented mainly on code simplification and spotting redundancies.

On the other hand, there are some metrics that evaluate how effectively trained models predict the acceptance of PR or time to close a PR; and these metrics can be compared across different programming languages. On top of that, the results from the chi-square test can also be compared. However, the caution is in order because the code quality for each language is evaluated differently, as discussed before.




Figure 6.30: Comparison of Cramer's V

As can be seen in Figure 6.30, in all studied languages, the presence of some issue have a negative effect on the PR acceptance (in terms of chi-square tests); however, for all of the languages, this effect is small ($\phi_c \approx 0.1$). The smallest effect was observed for Haskell programming language and the highest effect for C/C++. On the other hand, taking into account solely the PRs that only modified some source code of the primary language, the chi-square test indicates that the presence of issue and PR acceptance are independent in the case of the C/C++ and Kotlin. This is a possible threat to validity.

The effect of code quality on acceptance was also studied using classification algorithms. One of the metrics that were used to measure the performance of the classification models is the “area under the ROC curve” (AUC). When using this metric to evaluate models, the Haskell is once again the language when the code quality is least important (Figure 6.6). The average AUC for Haskell models is around 0.5 — the trained models are no better than random guessing. The models for the Kotlin are ranked with the highest AUC score and therefore are better in classification than models for other languages. Except for Haskell, the average AUC is over 0.5 but under 0.6 — these AUC scores are usually considered poor [21]. This indicates that code quality has only a small or no effect on the acceptance.

As can be seen, similar results were obtained for all of the languages. In all of the languages, the code quality impact is small (based on the chi-square tests and also based on the results from classification algorithms). There is no language that significantly differs from others.




Figure 6.31: AUC for different languages (ROC)

As discussed in the previous chapter, there seems to be no connection between the code quality and the time it takes to close pull requests (based on the trained regression models). The smallest MAE was scored by Kotlin models (around 26 days); on the other end of the spectrum are Python models with an average MAE equal to 78.9 days. The trained models are unusable, considering that most of the pull requests are closed within the first two weeks (83 % of Kotlin PRs and 76% of Python PRs).




Figure 6.32: Mean absolute error for prediction of time to close a PR

6.7 Threats to validity

The validity of my research is endangered by several things. At first, the selection of the projects is one of the factors that influence the outcomes of the research. This study is focused primarily on popular projects. The rationale behind the project selection is explained in the own dedicated subchapter. It is possible that projects selected using different metrics can yield varying results.

Another possible threat to validity is the selection of pull requests. It is usually not doable to examine all the pull requests of some project. For the projects with a huge number of PRs, the time and computational resources are the limiting factors. Moreover, to examine the rejected pull requests, the forked repository with the required commits needs to be available. This is not always the case. Sometimes the *force push* can also remove the commits from the accepted pull requests. It is also important to note that linting of some pull requests resulted

in an error in the linter, and therefore these PRs were skipped. Pull requests were also skipped if the linting time exceeded the limit (that was set to 1000 seconds) — the PRs that modified a huge number of files were ignored.

Another problem is that pull requests can be merged manually outside GitHub. These pull requests are not recognized as accepted. The projects were selected so that GitHub is the primary way to merge PRs. However, there still can be some PRs merged using alternative methods.

Furthermore, different methods can be used to measure the quality of pull requests. For each programming language, there exist several linters that are focused on a different set of issues, and they can also use different algorithms to detect the same issue. Another possible threat are false positives from linters. The false-positive can arise due to the fact that the files were linted in the isolated environment, and this can introduce some issues (`import-error`, etc.). Some issues are also hard to detect; for instance, the issue can be specific to some particular context, and the linter does not have to take this context into account. The greatest difficulty with the quality evaluation is the fact that everyone has a unique personal perspective on code quality — code quality does not have a single definition.

The pull requests sometimes contain also files that are not written in the primary programming language of the project. The pull request then can be rejected because of these files.

Lastly, there are several factors that influence PR acceptance. Some of them were discussed in previous chapters (number of commits, submitter's status, etc.). The one factor that influences the acceptance is a number of lines that were changed [7]. The more lines are added/changed, the higher the probability that the pull request will be rejected, but the chance that some quality issue will be introduced is also higher. In this case, it is difficult to distinguish if the pull request was rejected because of the code quality or because the changes are too big.

7 TODO Conclusion

7.1 TODO Comparison with the previous results

7.2 TODO Future work

- try different linters for the same language
- do the analysis only with filtered PRs (only modified and only primary language)
 - more data needs to be collected (different dataset)

Appendix

Projects

Table 7.1: Python projects

Project	Stars	Analyzed PRs	Accepted	Rejected	Introduced issues	Fixed issues
pallets/flask	58380	500	76%	24%	2.82	1.88
rg3/youtube-dl	50768	808	49%	51%	5.34	2.11
psf/requests	47100	500	64%	36%	2.76	1.04
nvbn/thefuck	46148	268	47%	53%	6.61	0.98
scrapy/scrapy	43124	500	80%	20%	5.63	5.39
faif/python-patterns	31006	258	90%	10%	4.20	5.30
certbot/certbot	28785	500	81%	19%	3.62	1.60
openai/gym	26986	500	68%	32%	7.12	3.05
soimort/you-get	25437	487	56%	44%	7.77	1.91
explosion/spaCy	23007	500	91%	9%	3.68	3.79
pypa/pipenv	22785	500	86%	14%	5.62	1.78
keon/algorithms	20528	341	81%	19%	11.42	9.82
tornadoweb/tornado	20451	500	80%	20%	3.18	1.36
keras-team/keras	20384	398	53%	47%	4.88	3.49
celery/celery	18850	500	76%	24%	3.60	1.24
locustio/locust	18518	496	74%	26%	9.02	3.90
sanic-org/sanic	15958	500	81%	19%	4.86	1.95
spotify/luigi	15485	500	72%	28%	5.62	2.96
kivy/kivy	14471	500	89%	11%	2.50	1.18
powerline/powerline	13187	396	81%	19%	23.88	4.10

Issue categories

Model reliability

- for regression and classification
- check the Lenarduzzi paper
- table with recall, precision etc.

Table 7.2: Java projects

Project	Stars	Analyzed PRs	Accepted	Rejected	Introduced issues	Fixed issues
iluwatar/java-design-patterns	49353	258	47%	53%	47.50	56.09
TheAlgorithms/Java	45228	430	36%	64%	22.60	96.09
ReactiveX/RxJava	40697	722	45%	55%	64.45	9.55
apache/dubbo	37035	500	85%	15%	23.98	44.06
PhilJay/MPAndroidChart	34862	251	66%	34%	18.04	8.97
square/retrofit	32964	493	62%	38%	15.80	4.66
bumptech/glide	32402	343	75%	25%	15.06	2.64
netty/netty	28942	500	83%	17%	15.75	12.72
apolloconfig/apollo	26588	500	86%	14%	29.91	2.23
JakeWharton/butterknife	25699	243	88%	12%	22.94	5.96
alibaba/druid	25294	500	94%	6%	22.05	20.24
alibaba/fastjson	24218	443	84%	16%	19.45	9.24
Netflix/Hystrix	22506	500	92%	8%	37.28	10.98
libgdx/libgdx	19848	500	80%	20%	10.99	5.84
google/ExoPlayer	19119	500	64%	36%	62.70	49.39
mybatis/mybatis-3	16982	500	75%	25%	22.13	7.71
arduino/Arduino	12729	500	58%	42%	1702.66	100.61
apache/hadoop	12429	500	78%	22%	20.44	25.30
pinpoint-pinpoint	12107	500	86%	14%	54.86	31.20
apm/pinpoint						
android-async-http	10654	204	69%	31%	10.41	5.33
http/android-async-http						

Table 7.3: Kotlin projects

Project	Stars	Analyzed PRs	Accepted	Rejected	Introduced issues	Fixed issues
square/okhttp	41886	430	84%	16%	31.24	26.01
JetBrains/kotlin	40892	500	43%	57%	5.83	10.22
square/leakcanary	27408	300	88%	12%	40.75	6.66
tachiyomior/tachiyomi	15623	252	63%	37%	11.07	1.41
android/compose-samples	11220	300	87%	13%	0.06	0.00
Kotlin/kotlinx.coroutines	10586	398	64%	36%	11.38	10.02
ktorio/ktor	9559	480	66%	34%	4.64	2.12
mozilla-mobile/fenix	5694	500	79%	21%	0.15	0.05
arrow-kt/arrow	4918	442	77%	23%	29.69	2.00
cashapp/sqlelight	4497	392	91%	9%	12.70	3.46
intellij-rust/intellij-rust	3873	500	98%	2%	5.94	20.64
gradle/kotlin-dsl-samples	3526	278	91%	9%	8.02	2.62
kotest/kotest	3378	341	81%	19%	33.86	9.02
square/kotlinpoet	3103	401	93%	7%	16.96	1.68
edvin/tornadofx	2457	254	83%	17%	5.57	1.15
Kotlin/dokka	2424	446	77%	23%	7.90	10.62
mozilla-mobile/android-components	1913	500	79%	21%	0.46	0.11
DroidKaigi/conference-app-2018	1222	292	97%	3%	3.98	8.42
JetBrains/kotlin-wrappers	1000	294	97%	3%	0.84	0.20
wordpress-mobile/AztecEditor-Android	356	214	97%	3%	7.13	1.95

Table 7.4: Haskell projects

Project	Stars	Analyzed PRs	Accepted	Rejected	Introduced issues	Fixed issues
jgm/pandoc	15134	361	44%	56%	1.20	0.59
purescript/purescript	7632	456	80%	20%	0.48	0.15
carp-lang/Carp	4476	340	93%	7%	0.62	0.17
unisonweb/unison	4314	436	94%	6%	0.64	0.49
input-output-hk/cardano-sl	3752	488	86%	14%	1.99	2.43
commercialhaskell/stack	3244	519	73%	27%	0.33	0.20
haskell/haskell-ide-engine	2428	405	94%	6%	0.55	0.39
wireapp/wire-server	2425	223	78%	22%	1.51	0.53
yesodweb/yesod	2410	336	62%	38%	0.23	0.04
simonmichael/hledger	1962	262	84%	16%	0.62	0.44
agda/agda	1780	298	84%	16%	0.80	0.56
diku-dk/futhark	1737	287	95%	5%	0.24	0.10
ekmett/lens	1617	200	67%	33%	0.11	0.00
ndmitchell/hlint	1261	301	98%	2%	0.04	0.00
haskell-servant/servant	1234	217	51%	49%	0.43	0.09
haskell/cabal	1178	847	63%	37%	0.22	0.14
haskell/aeson	1099	246	82%	18%	0.30	0.06
clash-lang/clash-compiler	1092	241	88%	12%	1.00	0.36
ucsd-progsys/liquidhaskell	687	253	77%	23%	2.37	0.85
yesodweb/wai	579	233	63%	37%	0.27	0.11

Table 7.5: C/C++ projects

Project	Stars	Analyzed PRs	Accepted	Rejected	Introduced issues	Fixed issues
microsoft/terminal	82226	500	81%	19%	0.00	0.02
nlohmann/json	29526	353	80%	20%	0.02	0.00
nothings/stb	18491	368	62%	38%	0.08	0.00
mpv-player/mpv	18296	500	53%	47%	0.25	0.38
simdjson/simdjson	15446	413	72%	28%	0.67	0.24
micropython/micropython	13951	500	43%	57%	0.14	0.08
hashcat/hashcat	12288	500	92%	8%	0.72	0.14
Tencent/rapidjson	11910	392	91%	9%	0.10	0.00
davisking/dlib	11058	275	79%	21%	0.04	0.03
reactos/reactos	10812	500	71%	29%	0.15	0.12
pybind/pybind11	10799	500	75%	25%	0.06	0.04
libevent/libevent	8693	293	58%	42%	0.34	0.79
irungentoo/toxcore	8667	500	90%	10%	0.85	0.75
libgit2/libgit2	8356	500	92%	8%	0.33	0.22
zeromq/libzmq	7678	500	96%	4%	0.52	0.33
Z3Prover/z3	7287	500	92%	8%	0.01	0.00
nodemcu/nodemcu-firmware	7028	500	78%	22%	1.48	0.93
minetest/minetest	6956	500	60%	40%	1.75	0.10
microsoft/cpprestsdk	6815	180	67%	33%	0.12	0.01
sass/libsass	4273	500	87%	13%	0.00	0.00

Table 7.6: Pylint issue categories

Category	Introduced in total	#PRs which introduced	Fixed in total	#PRs which fixed
warning	48931	3350	36865	1910
error	24657	2540	18841	1310
convention	91770	4683	76324	2447
refactor	16317	2543	14964	1483
info	2	1	2	1

Table 7.7: PMD issue categories

Category	Introduced in total	#PRs which introduced	Fixed in total	#PRs which fixed
Code Style	1341883	5780	522255	3387
Design	189046	4689	212123	2617
Documentation	343493	4524	222088	2032
Error Prone	157278	3435	60056	2062
Multithreading	9475	938	7818	825
Best Practices	177369	3822	128873	2133
Performance	15125	1512	18002	1092

ROC curves

Bibliography

1. YU, Y.; WANG, H.; FILKOV, V.; DEVANBU, P.; VASILESCU, B. Wait for It: Determinants of Pull Request Evaluation Latency on GitHub. In: *Working Conference on Mining Software Repositories*. IEEE, 2015, pp. 367–371. Available from DOI: 10.1109/MSR.2015.42.
2. TRAUTSCH, A.; HERBOLD, S.; GRABOWSKI, J. A longitudinal study of static analysis warning evolution and the effects of PMD on software quality in Apache open source projects. *Empirical Software Engineering*. 2020, vol. 25, pp. 5137–5192. Available from DOI: 10.1007/s10664-020-09880-1.
3. GOUSIOS, G.; ZAIDMAN, A.; STOREY, M.; DEURSEN, A. Work Practices and Challenges in Pull-Based Development: The Integrator’s Perspective. In: *International Conference on Software Engineering*. IEEE, 2015, vol. 1, pp. 358–368. Available from DOI: 10.1109/ICSE.2015.55.
4. KONONENKO, O.; ROSE, T.; BAYSAL, O.; GODFREY, M.; THEISEN, D.; WATER, B. Studying Pull Request Merges: A Case Study of Shopify’s Active Merchant. In: *International Conference on Software Engineering*. ACM, 2018, pp. 124–133. Available from DOI: 10.1145/3183519.3183542.
5. LI, Z.; YU, Y.; WANG, T.; YIN, Gang; LI, Shanshan; WANG, Huaimin. Are You Still Working on This An Empirical Study on Pull Request Abandonment. *Transactions on Software Engineering*. 2021, pp. 1–1. Available from DOI: 10.1109/TSE.2021.3053403.
6. ALAMI, A.; COHN, L.; WĄSOWSKI, A. How Do FOSS Communities Decide to Accept Pull Requests? In: *Proceedings of the Evaluation and Assessment in Software Engineering*. ACM, 2020, pp. 220–229. Available from DOI: 10.1145/3383219.3383242.
7. TSAY, J.; DABBISH, L.; HERBSLEB, J. Influence of Social and Technical Factors for Evaluating Contribution in GitHub. In: *International Conference on Software Engineering*. ACM, 2014, pp. 356–366. Available from DOI: 10.1145/2568225.2568315.

8. SOARES, D.; DE LIMA, M.; MURTA, L.; PLASTINO, A. Acceptance Factors of Pull Requests in Open-Source Projects. In: *Symposium on Applied Computing*. ACM, 2015, pp. 1541–1546. Available from DOI: 10.1145/2695664.2695856.
9. DEY, T.; MOCKUS, A. Effect of Technical and Social Factors on Pull Request Quality for the NPM Ecosystem. In: *International Symposium on Empirical Software Engineering and Measurement*. ACM, 2020. Available from DOI: 10.1145/3382494.3410685.
10. JOSH, J.; KOFINK, A.; MIDDLETON, J.; RAINEAR, C.; MURPHY-HILL, E.; PARNIN, C.; STALLINGS, J. Gender differences and bias in open source: Pull request acceptance of women versus men. *PeerJ Computer Science*. 2017, vol. 3. Available from DOI: 10.7717/peerj-cs.111.
11. IYER, R.; YUN, A.; NAGAPPAN, M.; HOEY, J. Effects of Personality Traits on Pull Request Acceptance. *Transactions on Software Engineering*. 2019. Available from DOI: 10.1109/TSE.2019.2960357.
12. GOLZADEH, M.; DECAN, A.; MENS, T. On the Effect of Discussions on Pull Request Decisions. In: *Belgium-Netherlands Software Evolution Workshop*. CEUR Workshop Proceedings, 2019. Available also from: <http://ceur-ws.org/Vol-2605/16.pdf>.
13. ZOU, W.; XUAN, J.; XIE, X.; CHEN, Z.; XU, B. How does code style inconsistency affect pull request integration? an exploratory study on 117 github projects. *Empirical Software Engineering*. 2019, vol. 24, pp. 3871–3903. Available from DOI: 10.1007/s10664-019-09720-x.
14. LENARDUZZI, V.; NIKKOLA, V.; SAARIMÄKI, N.; TAIBI, D. Does code quality affect pull request acceptance? An empirical study. *Journal of Systems and Software*. 2021, vol. 171, pp. 110806. Available from DOI: 10.1016/j.jss.2020.110806.
15. CHEN, D.; STOLEE, K.; MENZIES, T. Replication Can Improve Prior Results: A GitHub Study of Pull Request Acceptance. In: *International Conference on Program Comprehension*. IEEE, 2019, pp. 179–190. Available from DOI: 10.1109/ICPC.2019.00037.

16. ZHANG, X.; YU, Y.; GOUSIOS, G.; RASTOGI, A. Pull Request Decision Explained: An Empirical Overview. *Computing Research Repository*. 2021. Available from arXiv: 2105.13970.
17. GOUSIOS, G.; PINZGER, M.; DEURSEN, A. An Exploratory Study of the Pull-Based Software Development Model. In: *International Conference on Software Engineering*. ACM, 2014, pp. 345–355. Available from DOI: 10.1145/2568225.2568260.
18. DEY, T.; MOCKUS, A. Which Pull Requests Get Accepted and Why? A study of popular NPM Packages. *Computing Research Repository*. 2020. Available from arXiv: 2003.01153.
19. SOARES, D.; DE LIMA, M.; MURTA, L.; PLASTINO, A. Rejection Factors of Pull Requests Filed by Core Team Developers in Software Projects with High Acceptance Rates. In: *International Conference on Machine Learning and Applications*. IEEE, 2015, pp. 960–965. Available from DOI: 10.1109/ICMLA.2015.41.
20. AZEEM, I.; PENG, Q.; WANG, Q. Pull Request Prioritization Algorithm based on Acceptance and Response Probability. In: *International Conference on Software Quality*. IEEE, 2020, pp. 231–242. Available from DOI: 10.1109/QRS51102.2020.00041.
21. LEMESHOW, S.; STURDIVANT, R.; HOSMER, D. *Applied Logistic Regression*. Wiley, 2013. Wiley Series in Probability and Statistics. ISBN 9780470582473. Available from DOI: 10.1002/9781118548387.